

```
In [39]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, clas
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
In [2]: df = pd.read_csv("Social_Network_Ads.csv")
```

```
In [3]: df
```

```
Out[3]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...	...	...	...	...	...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

```
In [4]: df.shape
```

```
Out[4]: (400, 5)
```

```
In [5]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   User ID               400 non-null   int64  
1   Gender                400 non-null   object  
2   Age                   400 non-null   int64  
3   EstimatedSalary       400 non-null   int64  
4   Purchased             400 non-null   int64  
dtypes: int64(4), object(1)
memory usage: 15.8+ KB

```

In [6]: `df.describe()`

```

Out[6]:

```

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000


In [7]: `df.isna().sum()`

```


Out[7]:
User ID      0
Gender       0
Age          0
EstimatedSalary  0
Purchased    0
dtype: int64

```

In [9]: `histplot = sns.histplot(df['Age'], kde=True, bins=10, color='red', alpha=0.5)`
`for i in histplot.containers:`
 `histplot.bar_label(i,)`
`plt.show()`



```
In [10]: histplot = sns.histplot(df['EstimatedSalary'], kde=True, bins=10, color='
for i in histplot.containers:
    histplot.bar_label(i,)
plt.show()
```



```
In [11]: df["Gender"].value_counts()
```

```
Out[11]: Gender
Female    204
Male      196
Name: count, dtype: int64
```

```
In [12]: def gender_encoder(value):
        if (value == "Male"):
            return 1
        elif (value == "Female"):
            return 0
        else:
            return -1
```

```
In [13]: df["Gender"] = df["Gender"].apply(gender_encoder)
df
```

```
Out[13]:
```


	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	1	19	19000	0
1	15810944	1	35	20000	0
2	15668575	0	26	43000	0
3	15603246	0	27	57000	0
4	15804002	1	19	76000	0
...	...	...	...	...	...
395	15691863	0	46	41000	1
396	15706071	1	51	23000	1
397	15654296	0	50	20000	1
398	15755018	1	36	33000	0
399	15594041	0	49	36000	1

400 rows × 5 columns


```
In [14]: df["Purchased"].value_counts()
```

```
Out[14]: Purchased
0      257
1      143
Name: count, dtype: int64
```

```
In [15]: sns.pairplot(df, hue='Purchased', palette='bwr')
plt.show()
```



```
In [16]: sns.heatmap(df.corr(), annot=True)
plt.show()
```



```
In [17]: x = df[["User ID", "Gender", "Age", "EstimatedSalary"]]
         y = df["Purchased"]
```

```
In [18]: scaler = StandardScaler()
         x = scaler.fit_transform(x)
```

```
In [19]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
```

```
In [20]: x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

```
Out[20]: ((320, 4), (80, 4), (320,), (80,))
```

```
In [21]: model = LogisticRegression()
```

```
In [22]: model.fit(x_train, y_train)
```

```
Out[22]: LogisticRegression ⓘ ?
         Parameters
```

```
In [23]: y_pred = model.predict(x_test)
         y_pred
```

```
Out[23]: array([0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0,
               0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
               0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
               1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0])
```

```
In [24]: model.score(x_train,y_train)
```

```
Out[24]: 0.821875
```


```
In [25]: model.score(x,y)
```

```
Out[25]: 0.835
```

```
In [26]: cm = confusion_matrix(y_test, y_pred)
         print(cm)
```

```
[[50  2]
 [ 7 21]]
```

```
In [27]: disp=ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=model.classes_)
         disp.plot()
         plt.show()
```



```
In [28]: print(f"TN value is {cm[0][0]}")
         print(f"FP value is {cm[0][1]}")
         print(f"FN value is {cm[1][0]}")
         print(f"TP value is {cm[1][1]}")
```

```
TN value is 50
FP value is 2
FN value is 7
TP value is 21
```

```
In [40]: print(f"Accuracy score is {accuracy_score(y_test, y_pred)}")
print(f"Error rate is {1-accuracy_score(y_test, y_pred)}")
print(f"Precision score is {precision_score(y_test, y_pred)}")
print(f"Recall score is {recall_score(y_test, y_pred)}")
```

Accuracy score is 0.8875
Error rate is 0.11250000000000004
Precision score is 0.9130434782608695
Recall score is 0.75

```
In [32]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.96	0.92	52
1	0.91	0.75	0.82	28
accuracy			0.89	80
macro avg	0.90	0.86	0.87	80
weighted avg	0.89	0.89	0.88	80

```
In [ ]:
```